David Quispe

AI Engineer | LLM Evaluation, AI Safety and Security

Peru | Open to remote work (UTC-5) and relocation | da.qc@outlook.com | davidquicast.github.io
linkedin.com/in/davidquicast | github.com/davidquicast | huggingface.co/davidquicast

Summary

Systems engineer working where large language models meet security. Hands-on experience fine-tuning LLaMA 3.1 (8B and 70B) with PEFT and quantization, building RAG systems and tool-calling AI agents, and leading a vulnerability intelligence program in a public-sector security team. Co-author of INCLUDE-2.0 (EMNLP 2026), a multilingual LLM evaluation benchmark, and contributor to multilingual AI safety evaluation with the Cohere Labs community. Builder of PERSONA.md, an open specification and CLI for defining AI agent personas. Spanish (native), English (C1).

Experience

Cybersecurity Technician, FIAP (Public Administration Internationalization Foundation) Sep 2024 – Jun 2025
Madrid, Spain (hybrid) | Spanish public foundation for international governance cooperation

- Initiated and led a Vulnerability Intelligence program to proactively identify and remediate internal cybersecurity risks.
- Built proofs of concept with Microsoft Copilot, open-source LLMs and agentic frameworks to automate Cyber Threat Intelligence (CTI).
- Supported asset inventory, security auditing and incident response with EDR solutions and CCN-CERT tools.
- Hardened network security by managing and optimizing Fortinet technologies, reducing exploitable vulnerabilities.
- Secured the Microsoft 365 environment (SharePoint, Exchange, Defender) by applying Zero Trust (ZTNA) principles.
- Contributed to information security policies and compliance with ENS, NIST and CCN-CERT standards.

AI/ML Engineer, PlawLabs Ltd. Jul 2024 – Sep 2024
London, UK (remote) | AI virtual assistant and custom software startup | 1 direct report

- Fine-tuned and optimized LLaMA 3.1 8B and 70B with quantization and PEFT, building custom adapters merged with base models.
- Developed and deployed a Retrieval-Augmented Generation (RAG) system with an embedding model and a vector database.
- Deployed AI agents with LangChain and Langroid using tool calling and task management to automate processes and integrate external applications.
- Evaluated open-source and proprietary LLMs for client tasks and improved output accuracy with few-shot and active prompting.

Full Stack Web Developer (Internship), HAMB Ingeniería y Construcción S.A.C. Jan 2022 – Sep 2022
Cusco, Peru (on-site) | 2 direct reports

- Led an interdisciplinary team to design, build and deploy the company web application (Laravel, React, MySQL, WordPress).
- Implemented SAP for inventory control, reducing inefficiencies by 30%.
- Improved the main web application, contributing to a 25% increase in the customer base, and raised sales team productivity by 20% through a CRM rollout and training.

Research and Publications

- **INCLUDE-2.0: Disentangling Knowledge Gaps in Multilingual LLM Evaluation.** EMNLP 2026 (Main Conference), co-author. Benchmark that separates language proficiency from region-specific factual knowledge. aclanthology.org/2026.emnlp-main.1344
- **CulturalRiddles: A Multicultural Benchmark of Riddles.** arXiv:2610.29968 (2026), Cohere Labs Community, co-author (Spanish, Peru). Tests culturally grounded reasoning across languages. arxiv.org/abs/2610.29968
- **Spanish NLP datasets.** KUNTUR and WikiHowES datasets credited in the #Somos600M paper; WikiHowES listed among the training datasets of RigoChat 2.

Selected Projects

- **Personaxis and PERSONA.md** (2026 – present). AI persona infrastructure: an open specification and CLI to define, validate, lint, diff and compile AI agent personas, plus a registry and deployment platform. personaxis.com

- **Sorry-Bench Multilingual** (Cohere Labs collaboration). Translation pipeline and released multilingual dataset for evaluating LLM safety refusal behavior across languages.
- **Open Deep Research for Vulnerability Intelligence.** Agentic security research system (smolagents) whose code agents query NIST NVD, CISA KEV, EPSS and Shodan to produce vulnerability intelligence reports; companion vulnerability reasoning dataset on Hugging Face.
- **Kuntur.** Peruvian legal-domain LLM in Spanish and Quechua, fine-tuned from Gemma 2B and published on Hugging Face.
- **Open-source contributions.** Hugging Face smol-course (Spanish translation, fine-tuning and dataset notebooks), Argilla Synthetic Data Generator (Docker packaging, Ollama local inference), Hugging Face Gemma Recipes (fine-tuning recipe).

Skills

- **AI and ML:** LLM fine-tuning (PEFT adapters, quantization), RAG, AI agents and tool calling, LLM evaluation and benchmarks, prompt engineering, PyTorch, TensorFlow, Keras, Hugging Face, LangChain, smolagents, CrewAI.
- **Vector databases:** Chroma, Qdrant, DataStax.
- **Security:** vulnerability management and intelligence, threat intelligence, incident response, EDR, Zero Trust, Microsoft Defender, Fortinet, Nessus, OWASP ZAP, Burp Suite, Wireshark, Autopsy, threat modeling (STRIDE).
- **Frameworks and standards:** MITRE ATT&CK, NIST, CIS Controls, OWASP Top 10, OWASP Top 10 for LLM and GenAI, ISO 27001, ISO 22301, ENS.
- **Programming:** Python, SQL, C#, C++, Java, JavaScript, TypeScript, Kotlin, Dart, Go, Shell. FastAPI, ASP.NET, Laravel, React, Flutter.
- **Data and cloud:** SQL Server, MySQL, Oracle, MongoDB, Redis; AWS, Microsoft Azure, Google Cloud; Docker, Git.

Education

Bachelor's Degree in Systems Engineering, Andean University of Cusco, Peru 2017 – 2022
Online Master's in Cybersecurity, CEUPE (European Center of Postgraduate Studies), Spain 2022 – 2023
Online Diploma in Cybersecurity, Universidad Mayor, Chile 2023
Higher Technical Degree in Network Systems Administration (dual), DIGITECH, Madrid, Spain 2023 – 2025

Certifications and Training

- ISO 27001 Internal Auditor and ISO 22301 Internal Auditor, TÜV NORD (2023).
- Expert in Transformer Models and Deep Learning for NLP, LenguajeNatural.AI (2024).
- Building RAG Agents with LLMs, NVIDIA (2025); The Reasoning Course and AI Agents Fundamentals, Hugging Face (2025).
- CyberOps Associate and Endpoint Security, Cisco Networking Academy (2023); Ethical Hacking and Computer Forensics, National Forensic Investigation Institute of Peru (2023).